

Time series anomaly detection

This tutorial is an introduction to time series anomaly detection using TensorFlow.

Anomaly detection is a problem with implementations in a wide variety of domains. It involves the identification of novel or unexpected observations or sequences within the data being captured.

In this tutorial different styles of LSTM models are trained using a series of data stock market prices. Reconstruction error between the original and predicted data is used to determine the anomalies. In the last step, anomalies are plotted over the original data to visualize the results.

Covid19

affected the markets so was not a surprise that the algorithms found most novel market behavior over the outbreaks waves and lockdowns.

This tutorial includes five sections:

1. Data preprocessing.
2. Training on LSTM neural network. LSTM stands for Long-Short Term Memory
3. Reconstruction of the data based on LSTM prediction functionality.
4. Mean Absolute Error (MAE) computation between original and reconstructed data.
5. Anomalies detection and plotting.

credit: TensorFlow.org for some of their work included in this tutorial

```
In [333]: # -*- coding: utf-8 -*-
        """
        Created on Sat Apr 10 18:22:15 2021

        @author: ion_g
        """

        #import libraries
        from tensorflow import keras
        from sklearn.preprocessing import StandardScaler
        import pandas as pd
        import numpy as np
        from sklearn.preprocessing import MinMaxScaler
        import matplotlib.pyplot as plt
        import tensorflow as tf
        import os
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, LSTM, Dropout, RepeatVector, TimeDistributed
        import IPython
        import IPython.display

        np.random.seed(1)
        tf.random.set_seed(1)
        clear_console = True #Clear console after each run?
```

Read stock prices from finance.yahoo.com

```
In [334]: %%time

df = pd.read_csv('https://query1.finance.yahoo.com/v7/finance/download/BKHY?peri

#Convert date from string to datetime64[ns]
df['Date'] = pd.to_datetime(df['Date'])
print(df['Date'].min(), df['Date'].max())

#Print some information about the input data
print('data shape:\n', df.shape)
print('start sequence data samples:\n', df[:5])
print('end sequence data samples:\n', df[-5:])
print('data frame columns:\n', df.info())
```

2008-02-01 00:00:00 2021-04-06 00:00:00
data shape:
(3317, 7)
start sequence data samples:

	Date	Open	High	Low	Close	Adj Close	Volume
0	2008-02-01	24.15	24.15	24.15	24.15	21.200817	100
1	2008-02-04	24.15	24.15	24.15	24.15	21.200817	0
2	2008-02-05	24.15	24.15	24.15	24.15	21.200817	0
3	2008-02-06	24.15	24.15	24.15	24.15	21.200817	0
4	2008-02-07	23.10	23.10	23.10	23.10	20.279041	200

end sequence data samples:

	Date	Open	High	Low	Close	Adj Close	Volume
3312	2021-03-30	37.209999	37.209999	37.209999	37.209999	37.209999	0
3313	2021-03-31	37.209999	37.209999	37.209999	37.209999	37.209999	0
3314	2021-04-01	37.209999	37.209999	37.209999	37.209999	37.209999	100
3315	2021-04-05	37.209999	37.209999	37.209999	37.209999	37.209999	0
3316	2021-04-06	37.209999	37.209999	37.209999	37.209999	37.209999	100

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3317 entries, 0 to 3316
Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	Date	3317 non-null	datetime64[ns]
1	Open	3317 non-null	float64
2	High	3317 non-null	float64
3	Low	3317 non-null	float64
4	Close	3317 non-null	float64
5	Adj Close	3317 non-null	float64
6	Volume	3317 non-null	int64

dtypes: datetime64[ns](1), float64(5), int64(1)
memory usage: 181.5 KB
data frame columns:
None
Wall time: 629 ms

Plot data

Get first impresion about the data that will be used in this tutorial

```
In [335]: fig, ax = plt.subplots(num=None, figsize=(14, 6), dpi=80, facecolor='w', edgecolor='k')
size = df.shape[0]
print(df['Open'])
ax.set_title("Stock prices")
ax.plot(df['Date'], df['Open'], '-', color='blue', animated = True, linewidth=1)
ax.plot(df['Date'], df['High'], '-', color='red', animated = True, linewidth=1)
ax.plot(df['Date'], df['Low'], '-', color='green', animated = True, linewidth=1)
ax.plot(df['Date'], df['Close'], '-', color='yellow', animated = True, linewidth=1)
plt.legend(['Open', 'High', 'Low', 'Close'])
plt.show()


fig, ax = plt.subplots(num=None, figsize=(14, 6), dpi=80, facecolor='w', edgecolor='k')
size = df.shape[0]
print(df['Open'])
ax.set_title("Daily transaction")
ax.plot(df['Date'], df['Volume'], '-', color='blue', animated = True, linewidth=1)
plt.legend(['Volume'])
plt.show()
print()
```

```
0      24.150000
1      24.150000
2      24.150000
3      24.150000
4      23.100000
```

...

```
3312    37.209999
3313    37.209999
3314    37.209999
3315    37.209999
3316    37.209999
```

Name: Open, Length: 3317, dtype: float64



```
0      24.150000
1      24.150000
2      24.150000
3      24.150000
4      23.100000
```

...

```
3312    37.209999
```


3313 37.209999

3314 37.209999

3315 37.209999

3316 37.209999

Name: Open, Length: 3317, dtype: float64



Split data into train and test samples

Using test data to validate the model could help even that the time series windowing will supply test target for each time sequence.

```
In [336]: train = df.loc[df['Date'] <= '2021-12-31']
test = df.loc[df['Date'] > '2021-12-31']
print(f'train shape {train.shape}, train shape {test.shape}')
print(f'train tail\n {train.tail()}\ntest head:\n {test.head()}')
if (clear_console):
    IPython.display.clear_output()
print()
```

1. Data preprocessing.

Normalize the data

It is important to scale features before training a neural network. Normalization is a common way of doing this scaling. "sklearn" package provides MinMaxScaler object to normalize data between 1 and -1 depending on Max and Min values for each feature.

```
In [337]: print('Transform data using scaler')
train_scaled = train.copy()
scalers={}
columns= ['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume']
for i in columns: # train.columns:
    scaler = MinMaxScaler(feature_range=(-1,1))
    s_s = scaler.fit_transform(train_scaled[i].values.reshape(-1,1))
    s_s=np.reshape(s_s,len(s_s))
    scalers[f'scaler_{i}'] = scaler
    train_scaled[i]=s_s
print('\nOriginal feature values:\n', train.head(10))
print('\nNormalized feature values:\n', train_scaled.head(10))

if (clear_console):
    IPython.display.clear_output()
print()
```

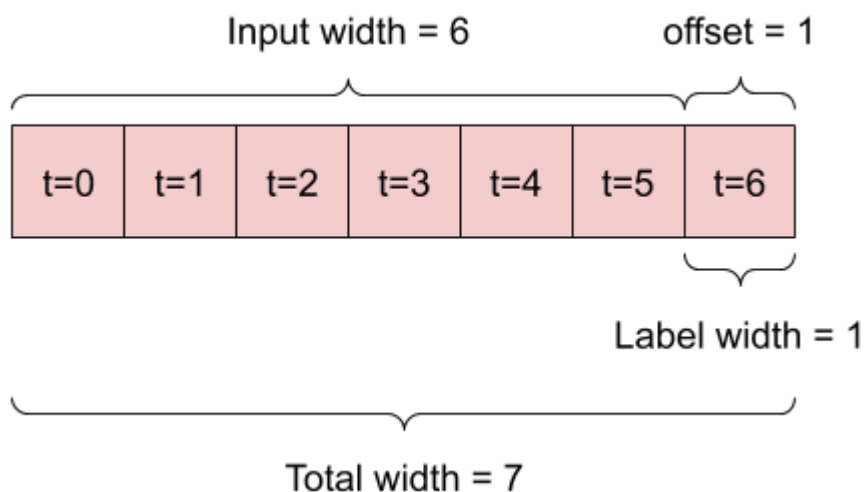
Data windowing

The models in this tutorial will make a set of predictions based on a window of consecutive samples from the data. The main features of the input windows are:

- The width (number of time steps) of the input and label windows
- The time offset between them.
- Which features are used as inputs, labels, or both.

In this tutorial I used a window of 30 day history (starting each day in the past) and one day test data.

For example, a model that makes a prediction of one step into the future, given six time steps of history would need a window like this



Credit to The TensorFlow Authors.

```

In [338]: TIME_STEPS=30
n_past = 30
n_future = 1
n_features = len(columns)
def split_series(series, n_past, n_future, dates ):
    #
    # n_past ==> no of past observations
    #
    # n_future ==> no of future observations
    #

    X, y, t = list(), list(), list()
    for window_start in range(len(series)):
        past_end = window_start + n_past
        future_end = past_end + n_future
        if future_end > len(series):
            break
        # slicing the past and future parts of the window
        past, future, date = series[window_start:past_end, :], series[past_end:future_end, :], dates[window_start:future_end]
        X.append(past)
        y.append(future)
        t.append(date)


    return np.array(X), np.array(y), t

def create_sequences(X, y, time_steps=TIME_STEPS):
    Xs, ys = [], []
    for i in range(len(X)-time_steps):
        Xs.append(X.iloc[i:(i+time_steps)].values)
        ys.append(y.iloc[i:(i+time_steps)].values)
    return np.array(Xs), np.array(ys)

X_train, y_train, time_series = split_series(train_scaled[columns].values, n_past, n_future, dates)
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], n_features))
y_train = y_train.reshape((y_train.shape[0], y_train.shape[1], n_features))
if (clear_console):
    IPython.display.clear_output()
print()

```

RNN model with one LSTM hidden layer



In [339]: # In[Build model]

```
model = Sequential()
model.add(LSTM(128, input_shape=(n_past, n_features)))
model.add(Dropout(rate=0.2))
model.add(RepeatVector(X_train.shape[1]))
model.add(LSTM(128, return_sequences=True))
model.add(Dropout(rate=0.2))
model.add(TimeDistributed(Dense(X_train.shape[2])))
model.compile(optimizer='adam', loss='mae')

model.summary()
if (clear_console):
    IPython.display.clear_output()
print()
```




```

In [340]: %%time
reduce_lr = tf.keras.callbacks.LearningRateScheduler(lambda x: 1e-3 * 0.90 ** x)
Epochs = 20

history = model.fit(X_train, y_train, epochs=Epochs, batch_size=32, validation_split=0.1,
                    callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', patience=10,
                                                             min_delta=0.001, mode='min'),
                               reduce_lr], shuffle=False)
IPython.display.clear_output()
print(f"model => loss: {history.history['loss'][-1]:.5f} val_loss: {history.history['val_loss'][-1]:.5f}")
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title("Model Loss")
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Valid'])
plt.show()

```

model => loss: 0.04580 val_loss: 0.04615



Wall time: 1min 45s

```
In [341]: %%time
# define Autoencoder model
#E2D2 128 Units
model128 = Sequential()
model128.add(LSTM(180, activation='relu', input_shape=(n_past, n_features), return_sequences=True))
model128.add(LSTM(32, activation='relu', return_sequences=False))
model128.add(RepeatVector((X_train.shape[1])))
model128.add(LSTM(32, activation='relu', return_sequences=True))
model128.add(LSTM(180, activation='relu', return_sequences=True))
model128.add(TimeDistributed(Dense(n_features)))

model128.compile(optimizer='adam', loss='mse')
model128.summary()
if (clear_console):
    IPython.display.clear_output()
print()
```

Wall time: 1.49 s

```

In [361]: %%time
history128 = model128.fit(X_train, y_train, epochs=Epochs, batch_size=32, validation_data=(X_val, y_val),
                           callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss',
                                                                      mode='min', verbose=1)],
                           verbose=1)


history128 = model128.fit(X_train, y_train, epochs=Epochs, batch_size=32, validation_data=(X_val, y_val),
                           callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss',
                                                                      mode='min', verbose=1)],
                           verbose=1)

IPython.display.clear_output()

print(f"model128 => loss: {history128.history['loss'][-1]:.5f} val_loss: {history128.history['val_loss'][-1]:.5f}")
plt.plot(history128.history['loss'])
plt.plot(history128.history['val_loss'])
plt.title("Model 128 Loss")
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Valid'])
plt.show()
pred128_e1d1=model1_e1d1.predict(X_train)

```

model128 => loss: 0.00409 val_loss: 0.00994



Wall time: 1min 58s

```

In [343]: %%time
# Model Architecture
#
# E1D1: One hidden layer encoder, one hidden layer decoder
# n_features ==> no of features at each timestep in the data.
# n_past ==> time steps in the past used for each sample at a given time
# Sample = X[t]
# Past = X[t-1],X[t-2],X[t-3],X[t-4],X[t-5],X[t-6],X[t-7],X[t-8],X[t-9],X[t-10]

encoder_inputs = tf.keras.layers.Input(shape=(n_past, n_features))
encoder_l1 = tf.keras.layers.LSTM(10, return_state=True)
encoder_outputs1 = encoder_l1(encoder_inputs)

encoder_states1 = encoder_outputs1[1:]

decoder_inputs = tf.keras.layers.RepeatVector(n_future)(encoder_outputs1[0])

decoder_l1 = tf.keras.layers.LSTM(10, return_sequences=True)(decoder_inputs, initial_state=encoder_states1)
decoder_outputs1 = tf.keras.layers.TimeDistributed(tf.keras.layers.Dense(n_features))(decoder_l1)

model_e1d1 = tf.keras.models.Model(encoder_inputs, decoder_outputs1)
model_e1d1.compile(optimizer='adam', loss='mae')
model_e1d1.summary()
if (clear_console):
    IPython.display.clear_output()
print()

```

Wall time: 513 ms

```


In [344]: %%time
# In[fit_model_e1d1]
history_e1d1 = model_e1d1.fit(X_train, y_train, epochs=Epochs*2, batch_size=32, validation_data=(X_val, y_val),
                             callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0.001, patience=5)])

IPython.display.clear_output()

print(f"model_e1d1 => loss: {history_e1d1.history['loss'][-1]:.5f} val_loss: {history_e1d1.history['val_loss'][-1]:.5f}")
plt.plot(history_e1d1.history['loss'])
plt.plot(history_e1d1.history['val_loss'])
plt.title("E1D1 Model Loss")
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Valid'])
plt.show()
pred1_e1d1=model_e1d1.predict(X_train)

```

model_e1d1 => loss: 0.02898 val_loss: 0.05541



Wall time: 22.1 s

```

In [345]: %%time
# Model Architecture
#
# E2D2: Two hidden layers encoder, two hidden layers decoder
# Convolution units: 100, 50, 1, 50, 100
# Return sequence and returns state (memory)
# n_features ==> no of features at each timestep in the data.
# n_past ==> time steps in the past used for each sample at a given time
# Sample = X[t]
# Past = X[t-1],X[t-2],X[t-3],X[t-4],X[t-5],X[t-6],X[t-7],X[t-8],X[t-9],X[t-10]

encoder_inputs = tf.keras.layers.Input(shape=(n_past, n_features))
encoder_l1 = tf.keras.layers.LSTM(100,return_sequences = True, return_state=True)
encoder_outputs1 = encoder_l1(encoder_inputs)
encoder_states1 = encoder_outputs1[1:]
encoder_l2 = tf.keras.layers.LSTM(50, return_state=True)
encoder_outputs2 = encoder_l2(encoder_outputs1[0])
encoder_states2 = encoder_outputs2[1:]

#Repeater => transfer data from encoder (Encoder outputs) to decoder (decoder inputs)
decoder_inputs = tf.keras.layers.RepeatVector(n_future)(encoder_outputs2[0])

#Decoders
decoder_l1 = tf.keras.layers.LSTM(50, return_sequences=True)(decoder_inputs,initial_state=encoder_states1)
decoder_l2 = tf.keras.layers.LSTM(100, return_sequences=True)(decoder_l1,initial_state=encoder_states2)
decoder_outputs2 = tf.keras.layers.TimeDistributed(tf.keras.layers.Dense(n_features))(decoder_l2)

model_e2d2 = tf.keras.models.Model(encoder_inputs,decoder_outputs2)
model_e2d2.compile(optimizer='adam', loss='mae')

model_e2d2.summary()
if (clear_console):
    IPython.display.clear_output()
print()

```

Wall time: 1.01 s

```

In [346]: %%time
# In[fit_model_e2d2]
history_e2d2 = model_e2d2.fit(X_train, y_train, epochs=Epochs, batch_size=32, val
                             callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss',

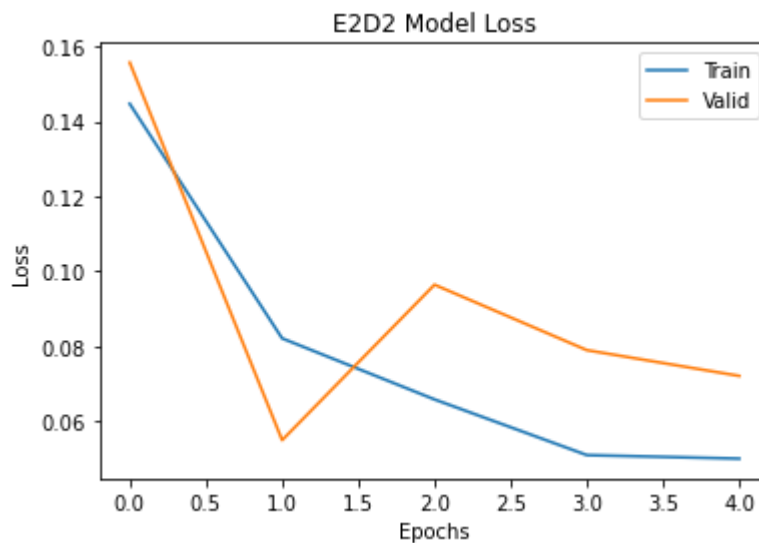
#history_e2d2=model_e1d1.fit(X_train, y_train, epochs=Epochs, batch_size=32,
                             # verbose=0,
#                             validation_split=0.1,
#                             callbacks=[reduce_lr])

IPython.display.clear_output()

print(f"model_e2d2 ==> loss: {history_e2d2.history['loss'][-1]:.5f}, val_loss: {h
plt.plot(history_e2d2.history['loss'])
plt.plot(history_e2d2.history['val_loss'])
plt.title("E2D2 Model Loss")
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Valid'])
plt.show()

```

model_e2d2 ==> loss: 0.04988, val_loss: 0.07199



Wall time: 45 s

```

In [347]: %%time
# Model Architecture
#
# E2D2: Two hidden layers encoder, two hidden layers decoder
# n_features ==> no of features at each timestep in the data.
# n_past ==> time steps in the past used for each sample at a given time
# Sample = X[t]
# Past = X[t-1],X[t-2],X[t-3],X[t-4],X[t-5],X[t-6],X[t-7],X[t-8],X[t-9],X[t-10]

encoder_inputs = tf.keras.layers.Input(shape=(n_past, n_features))
encoder_l1 = tf.keras.layers.LSTM(100, return_state=True, return_sequences = True)
encoder_outputs1 = encoder_l1(encoder_inputs)
encoder_l2 = tf.keras.layers.LSTM(50, return_state=True)
encoder_outputs2 = encoder_l2(encoder_outputs1[0])
encoder_states2 = encoder_outputs2[1:]

decoder_inputs = tf.keras.layers.RepeatVector(n_future)(encoder_outputs2[0])

decoder_l1 = tf.keras.layers.LSTM(50, return_sequences=True)
decoder_outputs1 = decoder_l1(decoder_inputs, initial_state = encoder_states2)
decoder_states2 = decoder_outputs1[1:]
decoder_l2 = tf.keras.layers.LSTM(100, return_sequences=True)
decoder_outputs2 = decoder_l2(decoder_outputs1)
decoder_states3 = decoder_outputs2[1:]
decoder_outputs3 = tf.keras.layers.TimeDistributed(tf.keras.layers.Dense(n_features))(decoder_outputs2)

model2_e2d2 = tf.keras.models.Model(encoder_inputs, decoder_outputs3)
model2_e2d2.compile(optimizer='adam', loss='mae')

model2_e2d2.summary()
if (clear_console):
    IPython.display.clear_output()
print()


```

Wall time: 5.3 s


```
In [348]: %%time
# In[fit model1_e2d2]
history2_e2d2=model2_e2d2.fit(X_train, y_train,epochs=Epochs,batch_size=32,
                              # verbose=0,
                              validation_split=0.1,
                              callbacks=[reduce_lr])

IPython.display.clear_output()
print(f"model2_e2d2 ==> loss: {history2_e2d2.history['loss'][-1]:.5f}, val_loss: {history2_e2d2.history['val_loss'][-1]:.5f}")
plt.plot(history2_e2d2.history['loss'])
plt.plot(history2_e2d2.history['val_loss'])
plt.title("E2D2 Model2 Loss")
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Valid'])
plt.show()
```

model2_e2d2 ==> loss: 0.01542, val_loss: 0.03209



Wall time: 1min 24s

Compare models

```
In [362]: #Compare models
print(f"model          ==> loss: {history.history['loss'][-1]:.5f}, val_loss: {history.history['val_loss'][-1]:.5f}")
print(f"model_e2d2     ==> loss: {history_e2d2.history['loss'][-1]:.5f}, val_loss: {history_e2d2.history['val_loss'][-1]:.5f}")
print(f"model_e1d1      ==> loss: {history_e1d1.history['loss'][-1]:.5f}, val_loss: {history_e1d1.history['val_loss'][-1]:.5f}")
print(f"model2_e2d2     ==> loss: {history2_e2d2.history['loss'][-1]:.5f}, val_loss: {history2_e2d2.history['val_loss'][-1]:.5f}")
print(f"model128        ==> loss: {history128.history['loss'][-1]:.5f}, val_loss: {history128.history['val_loss'][-1]:.5f}")


model          ==> loss: 0.04580, val_loss: 0.04615
model_e2d2     ==> loss: 0.04988, val_loss: 0.07199
model_e1d1      ==> loss: 0.02898, val_loss: 0.05541
model2_e2d2     ==> loss: 0.01542, val_loss: 0.03209
model128        ==> loss: 0.00409, val_loss: 0.00994
```

```

In [363]: %%time
# In[plot prediction]
pred_model=model.predict(X_train)
pred_e1d1=model_e1d1.predict(X_train)
pred_e2d2=model_e2d2.predict(X_train)
pred2_e2d2=model2_e2d2.predict(X_train)
pred128=model128.predict(X_train)

for i in range(len(columns)):
    fig, ax = plt.subplots(num=None, figsize=(14, 6), dpi=80, facecolor='w', edgecolor='k')
    ax.plot(time_series[-360:], X_train[-360:,29,i], '-', color='red', animated = True)
    ax.plot(time_series[-360:], pred_e1d1[-360:,0,i], '-', color='black', animated = True)
    ax.plot(time_series[-360:], pred_model[-360:,0,0], '-', color='yellow', animated = True)
    ax.plot(time_series[-360:], pred128[-360:,0,0], '-', color='green', animated = True)
    plt.title(f"{columns[i]} prediction")
    # plt.xlabel('Epochs')
    # plt.ylabel('Loss')
    plt.legend(['X_train', 'Predict e1d1', 'Predict simple', 'Predict 128'])
    plt.show()

```



```
In [364]: %%time
# Unscale data
# Convert prediction to original data scale

pred1_e2d2_unscaled=pred1_e2d2.copy()
pred1_e1d1_unscaled=pred1_e1d1.copy()
pred1_unscaled=pred1.copy()
pred2_e2d2_unscaled=pred2_e2d2.copy()
pred128_unscaled=pred128.copy()

for index,i in enumerate(columns):
    scaler = scalers[f'scaler_{i}']
    pred1_e2d2_unscaled[:, :, index]=scaler.inverse_transform(pred1_e2d2[:, :, index])
    pred1_e1d1_unscaled[:, :, index]=scaler.inverse_transform(pred1_e1d1[:, :, index])
    pred1_unscaled[:, :, index]=scaler.inverse_transform(pred1[:, :, index])
    pred2_e2d2_unscaled[:, :, index]=scaler.inverse_transform(pred2_e2d2[:, :, index])
    pred128_unscaled[:, :, index]=scaler.inverse_transform(pred128[:, :, index])
if (clear_console):
    IPython.display.clear_output()
print()
```

Wall time: 24 ms

```
In [365]: # In[calculate RMSE]
from math import sqrt
from sklearn.metrics import mean_squared_error

for i in range(len(columns)):
    rmse = sqrt(mean_squared_error(y_train[:, :, i], pred128[:, :, i]))
    print(f'{columns[i]:10s} Test RMSE: %.3f' % rmse)
```

Open	Test RMSE: 0.080
High	Test RMSE: 0.078
Low	Test RMSE: 0.083
Close	Test RMSE: 0.081
Adj Close	Test RMSE: 0.112
Volume	Test RMSE: 0.059

Anomaly detection

Anomaly is where reconstruction error is large. We can define this value beyond which we call anomaly. Let us look at MAE in training prediction

```


In [376]: %%time
trainMAE = np.mean(np.abs(pred128 - X_train), axis=1)
plt.title(f"Model128 mean error distribution")
plt.hist(trainMAE, bins=30)
plt.legend(columns)
histogram = np.histogram(trainMAE, bins=30, range=None, normed=None, weights=None)
print(histogram)
max_trainMAE = histogram[1][12] #Define max threshold.

anomaly_df = pd.DataFrame(pred128_unscaled[:,0,:],columns=columns)
for i in range(len(columns)):
    anomaly_df[f'MAE {columns[i]}'] = trainMAE[:,i]
    anomaly_df[f'anomaly {columns[i]}'] = anomaly_df[f'MAE {columns[i]}'] > max_t

if (clear_console):
    IPython.display.clear_output()
print()


```

Wall time: 124 ms



Plot anomalies

```
In [377]: %%time
for i in columns:
    fig, ax = plt.subplots(num=None, figsize=(14, 6), dpi=80, facecolor='w', edgecolor='k')
    ax.plot(time_series[-400:], anomaly_df[f'{i}'][-400:] / anomaly_df[f'anomaly {i}'], color='blue', marker='o')
    ax.plot(time_series[-400:], anomaly_df[f'{i}'][-400:], '-', color='green', linewidth=2)
    plt.title(f"{i} Anomalies {time_series[-361:-360]}-{time_series[-1:]}")
    plt.legend(['MAE', 'anomaly', '{1}'])
    plt.show()
```



In []: